

## 1. Data visualisation :

Our goal here is to visualize the database "heart\_failure\_clinical\_records\_dataset.csv" .

```
# reading the data file
data = pd.read_csv('heart_failure_clinical_records_dataset.csv')
# print the first 5 rows of the data
data.head()
```


|   | age  | anaemia | creatinine_phosphokinase | diabetes | ejection_fraction | high_blood_pressure | platelets | serum_creatinine | serum_sodium | sex | smoking | time |
|---|------|---------|--------------------------|----------|-------------------|---------------------|-----------|------------------|--------------|-----|---------|------|
| 0 | 75.0 | 0       | 582                      | 0        | 20                | 1                   | 265000.00 | 1.9              | 130          | 1   | 0       | 4    |
| 1 | 55.0 | 0       | 7861                     | 0        | 38                | 0                   | 263358.03 | 1.1              | 136          | 1   | 0       | 6    |
| 2 | 65.0 | 0       | 146                      | 0        | 20                | 0                   | 162000.00 | 1.3              | 129          | 1   | 1       | 7    |
| 3 | 50.0 | 1       | 111                      | 0        | 20                | 0                   | 210000.00 | 1.9              | 137          | 1   | 0       | 7    |
| 4 | 65.0 | 1       | 160                      | 1        | 20                | 0                   | 327000.00 | 2.7              | 116          | 0   | 0       | 8    |


```
] data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 299 entries, 0 to 298
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                   299 non-null    float64
1   anaemia                              299 non-null    int64
2   creatinine_phosphokinase             299 non-null    int64
3   diabetes                             299 non-null    int64
4   ejection_fraction                    299 non-null    int64
5   high_blood_pressure                  299 non-null    int64
6   platelets                            299 non-null    float64
7   serum_creatinine                     299 non-null    float64
8   serum_sodium                         299 non-null    int64
9   sex                                  299 non-null    int64
10  smoking                              299 non-null    int64
11  time                                 299 non-null    int64
12  DEATH_EVENT                          299 non-null    int64
dtypes: float64(3), int64(10)
memory usage: 30.5 KB
```

```
: data.isnull().sum()
```

```
: age 0
  anaemia 0
  creatinine_phosphokinase 0
  diabetes 0
  ejection_fraction 0
  high_blood_pressure 0
  platelets 0
  serum_creatinine 0
  serum_sodium 0
  sex 0
  smoking 0
  time 0
  DEATH_EVENT 0
  dtype: int64
```





## 2. Machine Learning applications before data Normalization:

### a. Model preparation:

```
data.columns
```

```
Index(['age', 'anaemia', 'creatinine_phosphokinase', 'diabetes',
      'ejection_fraction', 'high_blood_pressure', 'platelets',
      'serum_creatinine', 'serum_sodium', 'sex', 'smoking', 'time',
      'DEATH_EVENT'],
      dtype='object')
```

```
X_data = data[['age', 'anaemia', 'creatinine_phosphokinase', 'diabetes',
               'ejection_fraction', 'high_blood_pressure', 'platelets',
               'serum_creatinine', 'serum_sodium', 'sex', 'smoking', 'time']]
y_data = data['DEATH_EVENT']
```

```
X_train,X_test,y_train,y_test = train_test_split(X_data,y_data,test_size=0.3, stratify=y_data)
```

```
X_train,X_val,y_train,y_val = train_test_split(X_train,y_train,test_size=0.3, stratify=y_train)
```

b. Random forest :

```
clf = RandomForestClassifier(max_depth=2, random_state=0)
```

```
clf.fit(X_train,y_train)
```

```
RandomForestClassifier(max_depth=2, random_state=0)
```


```
#Predict the response for test dataset  
predictions = clf.predict(X_test)  
predictions
```

```
array([0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,  
       1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,  
       1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,  
       0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0,  
       0, 1], dtype=int64)
```

```
print(metrics.confusion_matrix(y_test,predictions))  
# using scikitlearn  
skplt.metrics.plot_confusion_matrix(y_test,predictions)
```

```
[[60  1]  
 [14 15]]
```

```
<AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>
```



```

acc = metrics.accuracy_score(y_test, predictions)
print("Accuracy = %.2f" %(acc))
f1 = metrics.f1_score(y_test, predictions)
print("F1 = %.2f" %(f1))
p = metrics.precision_score(y_test, predictions)
print("Precision = %.2f" %(p))
r = metrics.recall_score(y_test, predictions)
print("Recall = %.2f" %(r))
loss = metrics.log_loss(y_test, predictions)
print("log-loss = %.2f" %(loss))
auc = metrics.roc_auc_score(y_test, predictions)
print("ROC-AUC = %.2f" %(auc))

```

```

Accuracy = 0.83
F1 = 0.67
Precision = 0.94
Recall = 0.52
log-loss = 5.76
ROC-AUC = 0.75

```

c. Adaboost classifier:

```
ad=AdaBoostClassifier(base_estimator=clf)
```

```
ad.fit(X_test,y_test)
```

```
AdaBoostClassifier(base_estimator=RandomForestClassifier(max_depth=2,
                                                         random_state=0))
```

```
y_pred = ad.predict(X_val)
```

```

print(metrics.confusion_matrix(y_val,y_pred))
# using scikitlearn
skplt.metrics.plot_confusion_matrix(y_val,y_pred)


```

```

[[37  6]
 [ 7 13]]

```

```
<AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>
```



```

acc = metrics.accuracy_score(y_val,y_pred)
print("Accuracy = %.2f" %(acc))
f1 = metrics.f1_score(y_val,y_pred)
print("F1 = %.2f" %(f1))
p = metrics.precision_score(y_val,y_pred)
print("Precision = %.2f" %(p))
r = metrics.recall_score(y_val,y_pred)
print("Recall = %.2f" %(r))
loss = metrics.log_loss(y_val,y_pred)
print("log-loss = %.2f" %(loss))
auc = metrics.roc_auc_score(y_val,y_pred)
print("ROC-AUC = %.2f" %(auc))

```

Accuracy = 0.79  
 F1 = 0.67  
 Precision = 0.68  
 Recall = 0.65  
 log-loss = 7.13  
 ROC-AUC = 0.76

d. SVM :

```

SVM = svm.SVC(kernel='linear') # Linear Kernel

#Train the model using the training sets
SVM.fit(X_train, y_train)

#Predict the response for test dataset
y_pred = SVM.predict(X_test)

```

```

print(metrics.confusion_matrix(y_test, y_pred))
# using scikitlearn
skplt.metrics.plot_confusion_matrix(y_test, y_pred)


```

```

[[53  8]
 [10 19]]

```

<AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>



```

: # Model Accuracy: how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
# Model Precision: what percentage of positive tuples are labeled as such?
print("Precision:",metrics.precision_score(y_test, y_pred))

# Model Recall: what percentage of positive tuples are labelled as such?
print("Recall:",metrics.recall_score(y_test, y_pred))

print("F1 =" ,metrics.f1_score(y_test,y_pred))

```

```

Accuracy: 0.8
Precision: 0.7037037037037037
Recall: 0.6551724137931034
F1 = 0.6785714285714286

```

### e. Gradient Boosting Classifier:

```
gb_clf = GradientBoostingClassifier(n_estimators=100, learning_rate=0.5, max_depth=3, random_state=0).fit(X_train, y_train)
```

```

#Train the model using the training sets
gb_clf.fit(X_train, y_train)

```

```

#Predict the response for test dataset
y_pred = gb_clf.predict(X_test)

```

```

print(metrics.confusion_matrix(y_test, y_pred))
# using scikitlearn
skplt.metrics.plot_confusion_matrix(y_test, y_pred)

```


```

[[52  9]
 [ 8 21]]

```

```
<AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>
```

```
<AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>
```



```

# Model Accuracy: how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
# Model Precision: what percentage of positive tuples are labeled as such?
print("Precision:",metrics.precision_score(y_test, y_pred))

# Model Recall: what percentage of positive tuples are labelled as such?
print("Recall:",metrics.recall_score(y_test, y_pred))

print("F1 =" ,metrics.f1_score(y_test,y_pred))

```

```

Accuracy: 0.8111111111111111
Precision: 0.7
Recall: 0.7241379310344828
F1 = 0.711864406779661

```

f. Logistic Regression :

```

# instantiate the model (using the default parameters)
logreg = LogisticRegression()

# fit the model with data
logreg.fit(X_train,y_train)

#
y_pred=logreg.predict(X_test)

```

```

print(metrics.confusion_matrix(y_test, y_pred))
# using scikitlearn
skplt.metrics.plot_confusion_matrix(y_test, y_pred)

```

```


[[53  8]
 [ 9 20]]

```

```

<AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>

```





```

: # Model Accuracy: how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
# Model Precision: what percentage of positive tuples are labeled as such?
print("Precision:",metrics.precision_score(y_test, y_pred))

# Model Recall: what percentage of positive tuples are labelled as such?
print("Recall:",metrics.recall_score(y_test, y_pred))

print("F1 =" ,metrics.f1_score(y_test,y_pred))

```

Accuracy: 0.8111111111111111  
 Precision: 0.7142857142857143  
 Recall: 0.6896551724137931  
 F1 = 0.7017543859649122

### 3. Normalisation :

```

from sklearn.preprocessing import StandardScaler
import numpy as np

# 4 samples/observations and 2 variables/features
scaler = StandardScaler()
scaler.fit(X_data)
scaled_data_x=scaler.transform(X_data)
scaled_data_x=pd.DataFrame(scaled_data_x, columns=X_data.columns)

```

scaled\_data\_x

|     | age       | anaemia   | creatinine_phosphokinase | diabetes  | ejection_fraction | high_blood_pressure | platelets     | serum_creatinine | serum_sodium |
|-----|-----------|-----------|--------------------------|-----------|-------------------|---------------------|---------------|------------------|--------------|
| 0   | 1.192945  | -0.871105 | 0.000166                 | -0.847579 | -1.530560         | 1.359272            | 1.681648e-02  | 0.490057         | -1.504036    |
| 1   | -0.491279 | -0.871105 | 7.514640                 | -0.847579 | -0.007077         | -0.735688           | 7.535660e-09  | -0.284552        | -0.141976    |
| 2   | 0.350833  | -0.871105 | -0.449939                | -0.847579 | -1.530560         | -0.735688           | -1.038073e+00 | -0.090900        | -1.731046    |
| 3   | -0.912335 | 1.147968  | -0.486071                | -0.847579 | -1.530560         | -0.735688           | -5.464741e-01 | 0.490057         | 0.085034     |
| 4   | 0.350833  | 1.147968  | -0.435486                | 1.179830  | -1.530560         | -0.735688           | 6.517986e-01  | 1.264666         | -4.682176    |
| ... | ...       | ...       | ...                      | ...       | ...               | ...                 | ...           | ...              | ...          |
| 294 | 0.098199  | -0.871105 | -0.537688                | 1.179830  | -0.007077         | 1.359272            | -1.109765e+00 | -0.284552        | 1.447094     |
| 295 | -0.491279 | -0.871105 | 1.278215                 | -0.847579 | -0.007077         | -0.735688           | 6.802472e-02  | -0.187726        | 0.539054     |
| 296 | -1.333392 | -0.871105 | 1.525979                 | 1.179830  | 1.854958          | -0.735688           | 4.902082e+00  | -0.575031        | 0.312044     |
| 297 | -1.333392 | -0.871105 | 1.890398                 | -0.847579 | -0.007077         | -0.735688           | -1.263389e+00 | 0.005926         | 0.766064     |
| 298 | -0.912335 | -0.871105 | -0.398321                | -0.847579 | 0.585389          | -0.735688           | 1.348231e+00  | 0.199578         | -0.141976    |

299 rows x 12 columns

```
: X_data=scaled_data_x
```

```
: y_data
```

```
: 0      1  
   1      1  
   2      1  
   3      1  
   4      1
```

```
   ..
```

```
294     0
```

```
295     0
```

```
296     0
```

```
297     0
```

```
298     0
```

```
Name: DEATH_EVENT, Length: 299, dtype: int64
```

#### 4. Machine learning applications after Normalization:

##### a. Model preparation :

```
X_train,X_test,y_train,y_test = train_test_split(X_data,y_data,test_size=0.3, stratify=y_data)
```

```
X_train,X_val,y_train,y_val = train_test_split(X_train,y_train,test_size=0.3, stratify=y_train)
```

##### b. Random forest :

```
acc = metrics.accuracy_score(y_test, predictions)  
print("Accuracy = %.2f" %(acc))  
f1 = metrics.f1_score(y_test, predictions)  
print("F1 = %.2f" %(f1))  
p = metrics.precision_score(y_test, predictions)  
print("Precision = %.2f" %(p))  
r = metrics.recall_score(y_test, predictions)  
print("Recall = %.2f" %(r))  
loss = metrics.log_loss(y_test, predictions)  
print("log-loss = %.2f" %(loss))  
auc = metrics.roc_auc_score(y_test, predictions)  
print("ROC-AUC = %.2f" %(auc))
```

```
Accuracy = 0.83
```

```
F1 = 0.68
```

```
Precision = 0.89
```

```
Recall = 0.55
```

```
log-loss = 5.76
```

```
ROC-AUC = 0.76
```

Before :

Accuracy = 0.83  
F1 = 0.67  
Precision = 0.94  
Recall = 0.52  
log-loss = 5.76  
ROC-AUC = 0.75

Normalization with Random forest has a negative effect on the evaluation of the module

c. Adaboost classifier:

```
acc = metrics.accuracy_score(y_val,y_pred)
print("Accuracy = %.2f" %(acc))
f1 = metrics.f1_score(y_val,y_pred)
print("F1 = %.2f" %(f1))
p = metrics.precision_score(y_val,y_pred)
print("Precision = %.2f" %(p))
r = metrics.recall_score(y_val,y_pred)
print("Recall = %.2f" %(r))
loss = metrics.log_loss(y_val,y_pred)
print("log-loss = %.2f" %(loss))
auc = metrics.roc_auc_score(y_val,y_pred)
print("ROC-AUC = %.2f" %(auc))
```

Accuracy = 0.89  
F1 = 0.80  
Precision = 0.93  
Recall = 0.70  
log-loss = 3.84  
ROC-AUC = 0.84

Before :

Accuracy = 0.79  
F1 = 0.67  
Precision = 0.68  
Recall = 0.65  
log-loss = 7.13  
ROC-AUC = 0.76

Normalization with Adaboost classifier has a positive effect on the evaluation of the module and accuracy

d. SVM:

```

# Model Accuracy: how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
# Model Precision: what percentage of positive tuples are labeled as such?
print("Precision:",metrics.precision_score(y_test, y_pred))

# Model Recall: what percentage of positive tuples are labelled as such?
print("Recall:",metrics.recall_score(y_test, y_pred))

print("F1 =" ,metrics.f1_score(y_test,y_pred))

```

```

Accuracy: 0.8111111111111111
Precision: 0.7
Recall: 0.7241379310344828
F1 = 0.711864406779661

```

Before :

```

Accuracy: 0.8
Precision: 0.7037037037037037
Recall: 0.6551724137931034
F1 = 0.6785714285714286

```

Normalization with SVM has a positive effect on the evaluation of the module and accuracy

e. Gradient Boosting Classifier:

```

# Model Accuracy: how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
# Model Precision: what percentage of positive tuples are labeled as such?
print("Precision:",metrics.precision_score(y_test, y_pred))

# Model Recall: what percentage of positive tuples are labelled as such?
print("Recall:",metrics.recall_score(y_test, y_pred))

print("F1 =" ,metrics.f1_score(y_test,y_pred))

```

```

Accuracy: 0.8444444444444444
Precision: 0.8
Recall: 0.6896551724137931
F1 = 0.7407407407407408

```

Before :

```

Accuracy: 0.8111111111111111
Precision: 0.7
Recall: 0.7241379310344828
F1 = 0.711864406779661

```

Normalization with Gradient Boosting Classifier has a positive effect on the evaluation of the module and accuracy

f. Logistic Regression :

```
# Model Accuracy: how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
# Model Precision: what percentage of positive tuples are labeled as such?
print("Precision:",metrics.precision_score(y_test, y_pred))

# Model Recall: what percentage of positive tuples are labelled as such?
print("Recall:",metrics.recall_score(y_test, y_pred))

print("F1 =" ,metrics.f1_score(y_test,y_pred))
```

```
Accuracy: 0.8222222222222222
Precision: 0.7241379310344828
Recall: 0.7241379310344828
F1 = 0.7241379310344829
```

Before :

```
Accuracy: 0.8111111111111111
Precision: 0.7142857142857143
Recall: 0.6896551724137931
F1 = 0.7017543859649122
```


Normalization with Logistic Regression has a positive effect on the evaluation of the module and accuracy

5. Data balancing:

```
#preparation des donnees  
y = y_data  
#X (les autres) sont les variables qui précèdent la dernière  
X= X_data
```

```
sb.countplot(x=y,data=X)
```

```
<AxesSubplot:xlabel='DEATH_EVENT', ylabel='count'>
```



```
# Division de la bd  
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3, stratify=y)  
X_train,X_val,y_train,y_val = train_test_split(X_train,y_train,test_size=0.3, stratify=y_train)  
# Sur-échantillonnage  
rOs = RandomOverSampler()  
X_ro, y_ro = rOs.fit_resample(X_train, y_train)  
X_rtest, y_rtest = rOs.fit_resample(X_test, y_test)  
X_rval, y_rval = rOs.fit_resample(X_val, y_val)  
print(X_ro.shape)  
print(X_rtest.shape)  
print(X_rval.shape)
```


```
(198, 12)
```

```
(122, 12)
```

```
(86, 12)
```

```
sb.countplot(x=y_ro,data=X_ro)
```

```
<AxesSubplot:xlabel='DEATH_EVENT', ylabel='count'>
```



## 6. Machine learning applications after data balancing :

### a. model preparation :

```
: X_train, y_train=X_ro, y_ro  
X_test, y_test=X_rtest, y_rtest  
X_val, y_val=X_rval, y_rval
```

### b. Random forest :

```
acc = metrics.accuracy_score(y_test, predictions)  
print("Accuracy = %.2f" %(acc))  
f1 = metrics.f1_score(y_test, predictions)  
print("F1 = %.2f" %(f1))  
p = metrics.precision_score(y_test, predictions)  
print("Precision = %.2f" %(p))  
r = metrics.recall_score(y_test, predictions)  
print("Recall = %.2f" %(r))  
loss = metrics.log_loss(y_test, predictions)  
print("log-loss = %.2f" %(loss))  
auc = metrics.roc_auc_score(y_test, predictions)  
print("ROC-AUC = %.2f" %(auc))
```

```
Accuracy = 0.81  
F1 = 0.80  
Precision = 0.85  
Recall = 0.75  
log-loss = 6.51  
ROC-AUC = 0.81
```

Before :

Accuracy = 0.83  
F1 = 0.67  
Precision = 0.94  
Recall = 0.52  
log-loss = 5.76  
ROC-AUC = 0.75

Data balancing with Adaboost classifier has a positive effect on the evaluation of the module and accuracy

c. Adaboost classifier:

```
acc = metrics.accuracy_score(y_val,y_pred)
print("Accuracy = %.2f" %(acc))
f1 = metrics.f1_score(y_val,y_pred)
print("F1 = %.2f" %(f1))
p = metrics.precision_score(y_val,y_pred)
print("Precision = %.2f" %(p))
r = metrics.recall_score(y_val,y_pred)
print("Recall = %.2f" %(r))
loss = metrics.log_loss(y_val,y_pred)
print("log-loss = %.2f" %(loss))
auc = metrics.roc_auc_score(y_val,y_pred)
print("ROC-AUC = %.2f" %(auc))
```

Accuracy = 0.83  
F1 = 0.82  
Precision = 0.85  
Recall = 0.79  
log-loss = 6.02  
ROC-AUC = 0.83

Before :

Accuracy = 0.89  
F1 = 0.80  
Precision = 0.93  
Recall = 0.70  
log-loss = 3.84  
ROC-AUC = 0.84

Data balancing with Adaboost classifier has a negative effect to accuracy

Data balancing with Adaboost classifier has a positive effect on the evaluation of the module

d. SVM:



```

# Model Accuracy: how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
# Model Precision: what percentage of positive tuples are labeled as such?
print("Precision:",metrics.precision_score(y_test, y_pred))

# Model Recall: what percentage of positive tuples are labelled as such?
print("Recall:",metrics.recall_score(y_test, y_pred))

print("F1 =" ,metrics.f1_score(y_test,y_pred))

```

```

Accuracy: 0.7950819672131147
Precision: 0.7903225806451613
Recall: 0.8032786885245902
F1 = 0.7967479674796747

```

Before :

```

Accuracy: 0.8111111111111111
Precision: 0.7
Recall: 0.7241379310344828
F1 = 0.711864406779661

```

Data balancing with SVM has a negative effect to accuracy

Data balancing with SVM has a positive effect on the evaluation of the module

e. Gradient Boosting Classifier:

```

# Model Accuracy: how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
# Model Precision: what percentage of positive tuples are labeled as such?
print("Precision:",metrics.precision_score(y_test, y_pred))

# Model Recall: what percentage of positive tuples are labelled as such?
print("Recall:",metrics.recall_score(y_test, y_pred))

print("F1 =" ,metrics.f1_score(y_test,y_pred))

```

```

Accuracy: 0.7704918032786885
Precision: 0.8367346938775511
Recall: 0.6721311475409836
F1 = 0.7454545454545455

```

Before :

Accuracy: 0.8444444444444444  
Precision: 0.8  
Recall: 0.6896551724137931  
F1 = 0.7407407407407408

Data balancing with Gradient Boosting Classifier has a negative effect to accuracy

f. Logistic Regression :

```
: # Model Accuracy: how often is the classifier correct?  
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))  
# Model Precision: what percentage of positive tuples are labeled as such?  
print("Precision:",metrics.precision_score(y_test, y_pred))  
  
# Model Recall: what percentage of positive tuples are labelled as such?  
print("Recall:",metrics.recall_score(y_test, y_pred))  
  
print("F1 =" ,metrics.f1_score(y_test,y_pred))
```

Accuracy: 0.7622950819672131  
Precision: 0.7962962962962963  
Recall: 0.7049180327868853  
F1 = 0.7478260869565216

Before :

Accuracy: 0.8222222222222222  
Precision: 0.7241379310344828  
Recall: 0.7241379310344828  
F1 = 0.7241379310344829

Data balancing with Logistic Regression has a positive effect on the evaluation of the module

Data balancing with Gradient Boosting Classifier has a negative effect to accuracy

